

# Laptop-sized ML for Text, with Open Source

**Nick Burch**

Berlin Buzzwords 2023

# Nick Burch

Director of Engineering



**FILEC**

# Nick Burch

@Gagravarr

@nick@social.earth.li

# Code, Scripts, Slides

[github.com/Gagravarr/BBuzz23-LaptopML](https://github.com/Gagravarr/BBuzz23-LaptopML)

All code in slides is taken from here!



# Our talk was going to be...

- LLMs - Large Language Models
- Why you can't have one
- Simpler vector-space language models
- Word2vec, GloVe, ELMo and BERT
- How they differ from traditional text relevancy like TF-IDF
- How to run them with Open Source libraries
- How to tweak them to get better relevance or speed or memory
- What you can do even though you don't have a LLM

# But then...









# Llamas happened





# **LLaMA**

## **Large Language Model**

### **from Facebook (they're still around!)**

# LLaMA

[ai.facebook.com/blog/large-language-model-llama-meta-ai/](https://ai.facebook.com/blog/large-language-model-llama-meta-ai/)

<https://arxiv.org/abs/2302.13971v1>

**LLaMa changed my talk, and the ML-for-text  
world!**





# All images in this talk are from Stable Diffusion

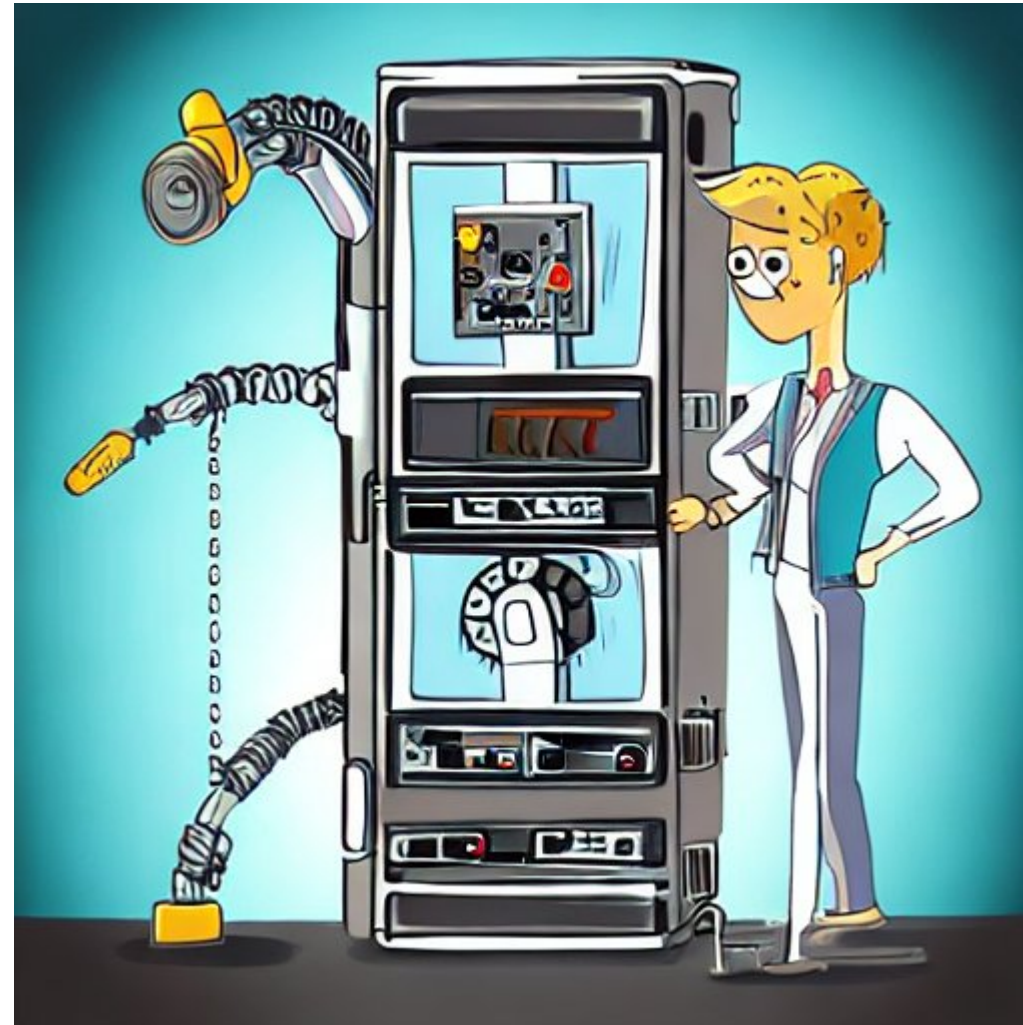
But this talk is all about text

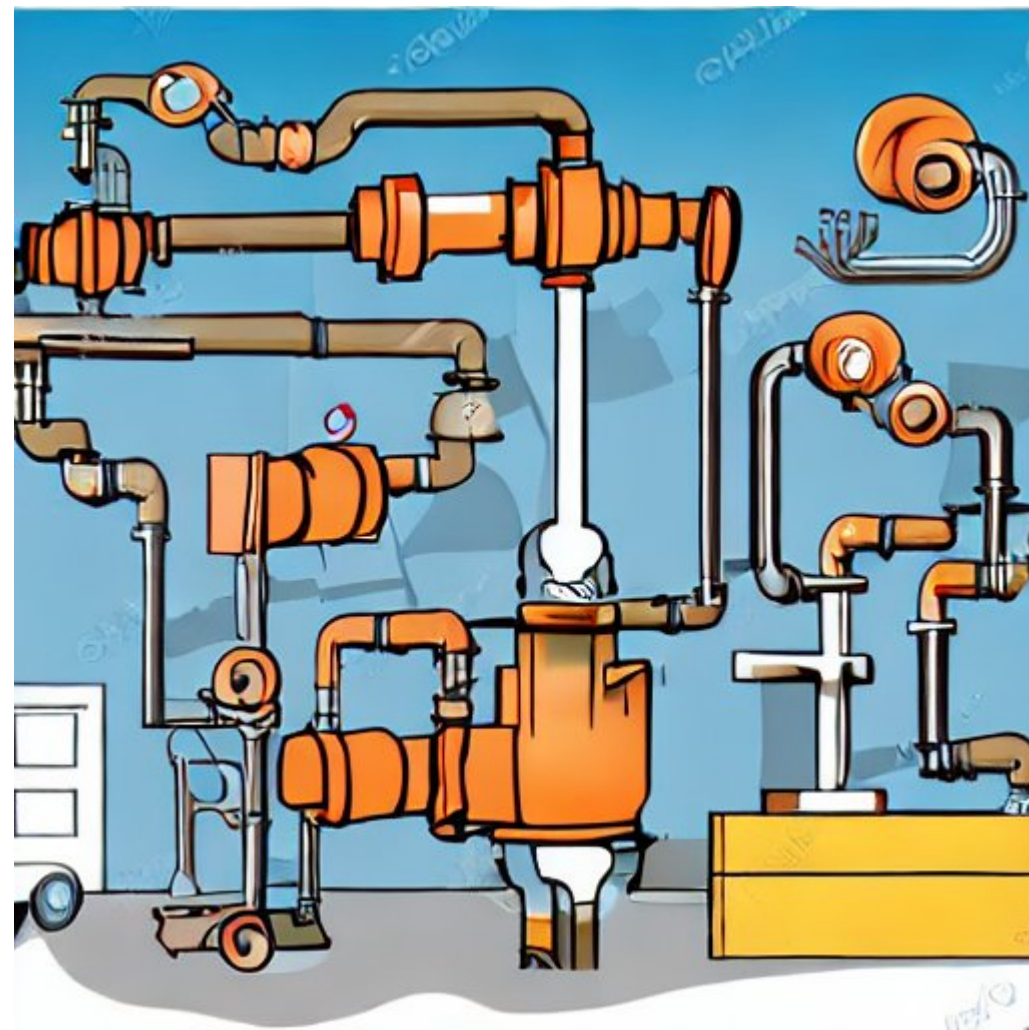
So we'll skip over all of that crazy-fun area of AI/ML!

[stablediffusionweb.com](https://stablediffusionweb.com)



**So what is an LLM, and/or a  
LLaMa?**





**LLM - Large Language Model**

**LLaMA - Facebook's (slightly open) LLM**

**LLaMa - Slightly open?**

**Code is open, model is not**

**But more on that later!**

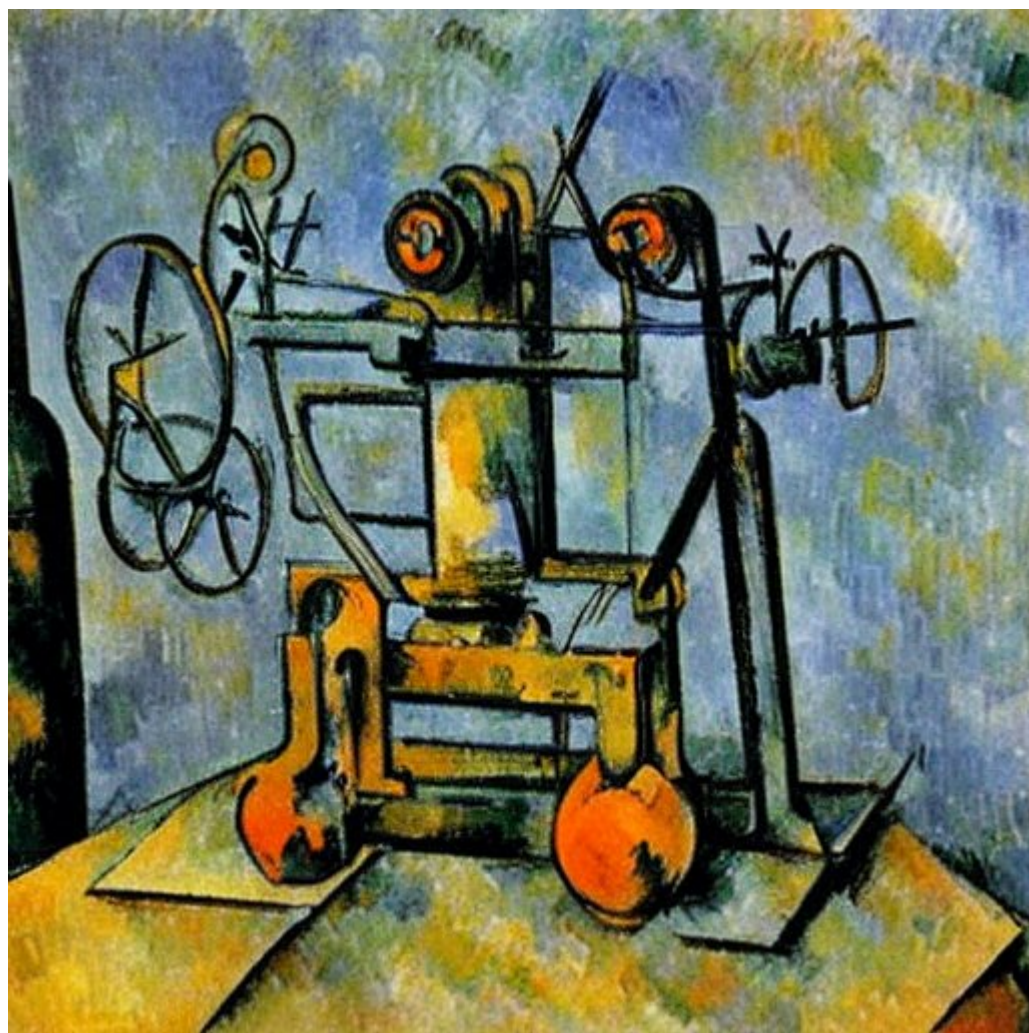
# LLM - Large Language Model?

A text-focused neural network with billions of parameters, trained on large amounts of text.

LLMs are deep-learning models, typically general purpose, which excel at a wide range of text-related tasks.

Some level of encoding on syntax and semantics of human language, and some level of general knowledge / facts.





Bigger models (both number of parameters and training corpus) will tend to show more knowledge and better semantics.



**Don't ignore the old stuff!**



# Everything you could do with BERT or ELMo, you can do with a LLM

- Embeddings? 
- Vector Search? 
- Similar Words? 
- Semantic Relationships? 

**So look back through the Buzzwords 2022 and  
2021 presentations**

**(They're all online on youtube)**

**If it mentioned BERT, you can try it with an  
LLM, and it'll (probably) be better!**



**But a LLM on a laptop?**

# Live demo time!

**llama.cpp, Facebook's LLaMa 7B (7 billion parameters)  
model, and a few of your questions**



# **Some ML + Information Retrieval terms to know**

# Tokenization

Breaking the input text into chunks, and possibly some simple transformations

*eg Hello there! Welcome to my talk.*

Could become

hello there welcome to my talk

Or... hello there [!] welcome to my talk [end]



# Embeddings

Both IR systems like Lucene, and ML techniques, need to work on numeric representations.

Lucene has a *term dictionary*, eg 1=hello 142=there

ML techniques have *embeddings*, a vector-space representation / projection,  
eg [0, 0, 0.9, 0, 0.1, 0.8]

# Embeddings


If this is all new to you, or you're a bit hazy on the details...

## What are embeddings

by *Vicki Boykis*



[github.com/veekaybee/what\\_are\\_embeddings](https://github.com/veekaybee/what_are_embeddings)



**What could the "old stuff" like  
BERT do?**

# Similar Words (semantically similar)

What other words are similar to X?

How similar are the words X and Y?

## Similar tokens to: linux

- Cosine sim=0.849: unix
- Cosine sim=0.793: open-source
- Cosine sim=0.778: kernel



Similar tokens to: raise

- Cosine sim=0.890: raising
- Cosine sim=0.871: pay
- Cosine sim=0.848: benefit



Difference between raise and risen is 56  
Difference between raise and above is 52  
Difference between raise and below is 54  
Difference between raise and shine is 30  
Difference between raise and linux is 11





# Word relationships

What is the equivalent of X to Y, for Z?

The analogy of X - Y for Z = ???

berlin - germany for paris = france

madrid - spain for lisbon = portugal

man - boy for woman = girl



# Warning - it can be wrong, it can be sexist!

The analogy of X - Y for Z = ???

spain - madrid for portugal = spain

doctor - man for nurse = woman

# Embeddings - sentence and/or document

Used in Vector Search (supported by Lucene, elastic etc)

- Transform the documents into the embedding vector space
- Transform the query into the embedding vector space
- Find documents similar to the query

Combination of embedding vector search and normal search can work well,  
see previous talks for details!



Rest to follow...

# Code, Scripts, Slides

[github.com/Gagravarr/BBuzz23-LaptopML](https://github.com/Gagravarr/BBuzz23-LaptopML)

All code in slides is taken from here!



